

# Insurance\_Data\_Set\_(SLR,\_MLR,\_PolyReg,\_Lasso,\_Ridge,\_ElasticNet)

May 21, 2024

## 0.1 1. Business Problem Understanding

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

## 0.2 2. Data Understanding

### 0.2.1 2.1 Data Collection

```
[2]: df = pd.read_excel('/content/insurance.xlsx')
df.head()
```

```
[2]:
```

	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   expenses    1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

# 1 Multiple Linear Regression

## 1.1 3. Data Preprocessing

### 1.1.1 3.1 EDA

```
[4]: df.info()
```

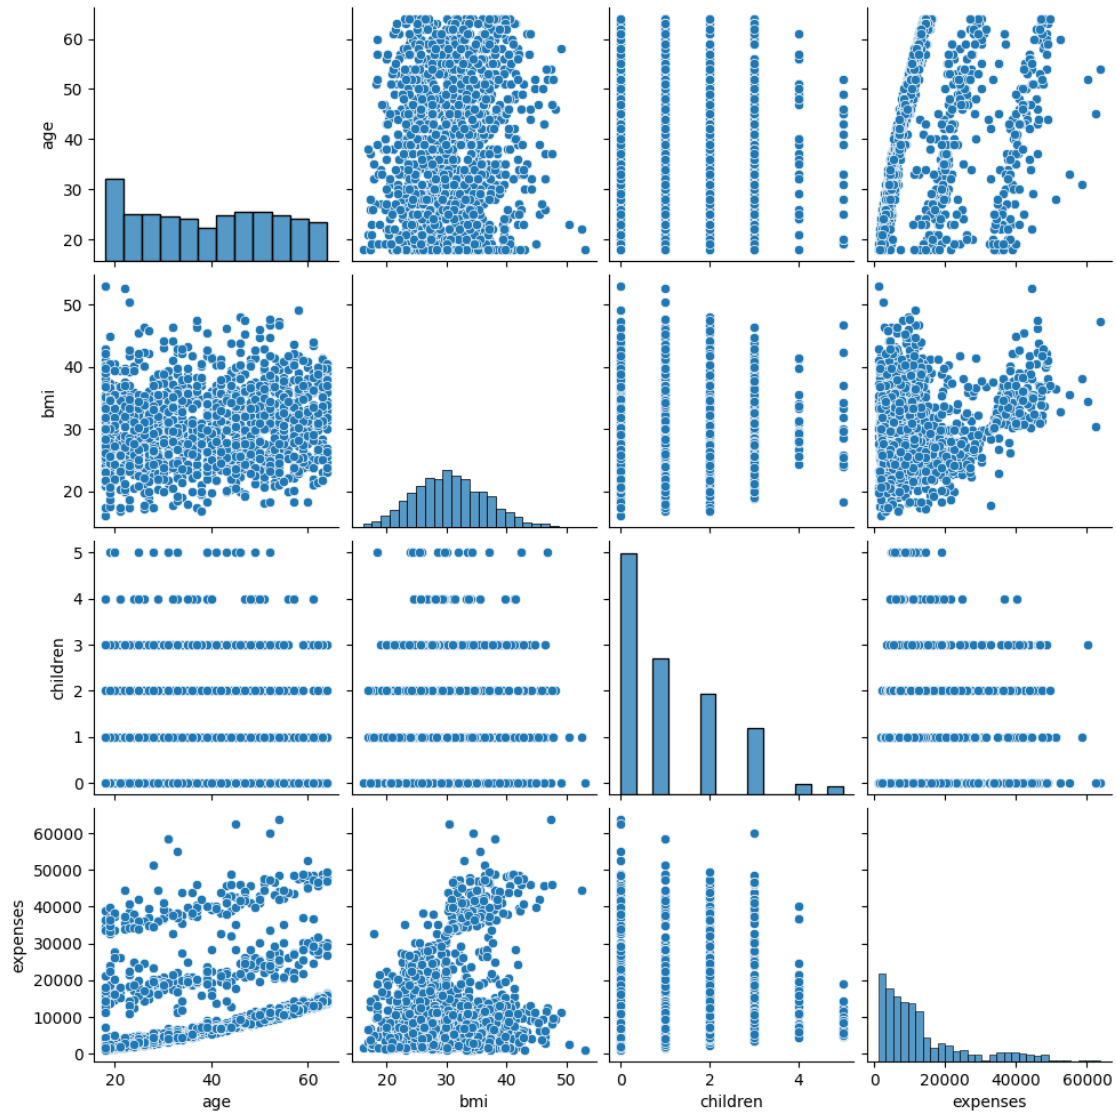
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   expenses    1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
[5]: df.describe()
```

```
[5]:
```

	age	bmi	children	expenses
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.665471	1.094918	13270.422414
std	14.049960	6.098382	1.205493	12110.011240
min	18.000000	16.000000	0.000000	1121.870000
25%	27.000000	26.300000	0.000000	4740.287500
50%	39.000000	30.400000	1.000000	9382.030000
75%	51.000000	34.700000	2.000000	16639.915000
max	64.000000	53.100000	5.000000	63770.430000

```
[6]: sns.pairplot(df)
plt.show()
```




```
[7]: df.corr(numeric_only=True)
```

```
[7]:
```

	age	bmi	children	expenses
age	1.000000	0.109341	0.042469	0.299008
bmi	0.109341	1.000000	0.012645	0.198576
children	0.042469	0.012645	1.000000	0.067998
expenses	0.299008	0.198576	0.067998	1.000000

```
[8]: sns.boxplot(data=df) # Checking the Outliers
plt.show()
```



```
[9]: df.skew(numeric_only=True)
```


```
[9]: age          0.055673
     bmi          0.284593
     children     0.938380
     expenses     1.515880
     dtype: float64
```

### 1.1.2 3.2 Data Cleaning

```
[10]: df.isnull().sum()
```

```
[10]: age          0
     sex          0
     bmi          0
     children     0
     smoker       0
     region       0
     expenses     0
     dtype: int64
```

```
[11]: # clearing the outliers
sns.boxplot(x=df['bmi'])
plt.show()
```



```
[12]: q1 = df['bmi'].quantile(0.25)
q3 = df['bmi'].quantile(0.75)
iqr = q3-q1
lower_limit = q1 - (iqr*1.5)
upper_limit = q3 + (iqr*1.5)
print(q1,q3,iqr,lower_limit,upper_limit)
```

26.3 34.7 8.400000000000002 13.699999999999998 47.300000000000004

```
[13]: df[df['bmi']>47.300000000000004] # Outliers of the Variable
```

```
[13]:
```

	age	sex	bmi	children	smoker	region	expenses
116	58	male	49.1	0	no	southeast	11381.33
286	46	female	48.1	2	no	northeast	9432.93
401	47	male	47.5	1	no	southeast	8083.92
543	54	female	47.4	0	yes	southeast	63770.43
847	23	male	50.4	1	no	southeast	2438.06

860	37	female	47.6	2	yes	southwest	46113.51
1047	22	male	52.6	1	yes	southeast	44501.40
1088	52	male	47.7	1	no	southeast	9748.91
1317	18	male	53.1	0	no	southeast	1163.46


```
[14]: df = df[(df['bmi'] > lower_limit) & (df['bmi'] < upper_limit)] # removing the
      ↪ outliers
      df
```

```
[14]:
```

	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86
...	...	...	...	...	...	...	...
1333	50	male	31.0	3	no	northwest	10600.55
1334	18	female	31.9	0	no	northeast	2205.98
1335	18	female	36.9	0	no	southeast	1629.83
1336	21	female	25.8	0	no	southwest	2007.95
1337	61	female	29.1	0	yes	northwest	29141.36

[1329 rows x 7 columns]

```
[15]: sns.boxplot(x=df['bmi'])
      plt.show()
```



```
[16]: df.drop('region',axis=1,inplace=True)
```

### 1.1.3 3.3 Data Wrangling

```
[17]: df['sex'].replace({'female':0,'male':1},inplace=True)
df['smoker'].replace({'no':0,'yes':1},inplace=True)
```

### 1.1.4 3.4 Train Test Split

```
[18]: x = df.drop('expenses',axis=1)
y = df['expenses']
```

```
[19]: from sklearn.model_selection import train_test_split as tts
x_train,x_test,y_train,y_test = tts(x,y, test_size = 0.2, random_state = 42)
```

```
[20]: x_test
```

```
[20]:
```

	age	sex	bmi	children	smoker
899	19	0	22.5	0	0
115	60	1	28.6	0	0
529	18	1	25.5	0	0

176	38	1	27.8	2	0
63	28	0	25.9	1	0
...	...	...	...	...	...
867	57	1	43.7	1	0
558	35	0	34.1	3	1
1028	54	1	31.6	0	0
585	33	0	28.3	1	0
1222	50	1	25.3	0	0

[266 rows x 5 columns]

## 1.2 4. Modelling and Evolution

```
[21]: # Modeling with default paramaters
from sklearn.linear_model import LinearRegression
mlr_model = LinearRegression()
mlr_model.fit(x_train,y_train)

# Prediction
train_predictions = mlr_model.predict(x_train)
test_predictions = mlr_model.predict(x_test)

# Evaluation
from sklearn.model_selection import cross_val_score
mlr_train = mlr_model.score(x_train,y_train)
mlr_test = mlr_model.score(x_test,y_test)
mlr_cv = cross_val_score(mlr_model,x,y,cv=5).mean()

print('Train R2: ',mlr_train)
print('Test R2: ',mlr_test)
print("CV Score: ",mlr_cv)
```

Train R2: 0.7436049992264148

Test R2: 0.768277484877281


CV Score: 0.7458418619962311

## 1.3 Check for Assumptions

### 1.3.1 1.Linearity of Errors


```
[22]: test_res = y_test - test_predictions
plt.scatter(y_test,test_res)
plt.xlabel("Observed Values")
plt.ylabel("Fitted values")
plt.show()
```






### 1.3.2 2.Normality of Errors

```
[23]: sns.displot(test_res,kde=True)  
plt.show()
```



### 1.3.3 3.Equal variance of errors (Homoscedasticity)

```
[24]: plt.scatter(test_predictions, test_res, c="r")  
plt.axhline(y=0, color='blue')  
plt.xlabel("Fitted values")  
plt.ylabel("Residuals")  
plt.show()
```



#### 1.3.4 4.Variables significance

```
[25]: import statsmodels.formula.api as smf
mod=smf.ols("y~x",data=df).fit()
mod.summary()
```

[25]:

Dep. Variable:	y	R-squared:	0.749
Model:	OLS	Adj. R-squared:	0.748
Method:	Least Squares	F-statistic:	788.5
Date:	Sat, 04 May 2024	Prob (F-statistic):	0.00
Time:	15:16:32	Log-Likelihood:	-13450.
No. Observations:	1329	AIC:	2.691e+04
Df Residuals:	1323	BIC:	2.694e+04
Df Model:	5		
Covariance Type:	nonrobust		

	coef	std err	t	P>  t	[0.025	0.975]
<b>Intercept</b>	-1.21e+04	964.025	-12.554	0.000	-1.4e+04	-1.02e+04
<b>x[0]</b>	257.0460	11.864	21.666	0.000	233.771	280.321
<b>x[1]</b>	-36.8547	332.019	-0.111	0.912	-688.195	614.486
<b>x[2]</b>	324.7149	28.140	11.539	0.000	269.511	379.919
<b>x[3]</b>	477.1451	137.039	3.482	0.001	208.307	745.983
<b>x[4]</b>	2.362e+04	411.630	57.377	0.000	2.28e+04	2.44e+04
<b>Omnibus:</b>		304.856	<b>Durbin-Watson:</b>		2.087	
<b>Prob(Omnibus):</b>		0.000	<b>Jarque-Bera (JB):</b>		730.634	
<b>Skew:</b>		1.236	<b>Prob(JB):</b>		2.21e-159	
<b>Kurtosis:</b>		5.661	<b>Cond. No.</b>		299.	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## 2 Polynomial Regression

### 2.1 3. Preprocessing

```
[26]: from sklearn.preprocessing import PolynomialFeatures
Poly_converter= PolynomialFeatures(degree=2,include_bias=False)
Poly_x = pd.DataFrame(Poly_converter.fit_transform(x))
```

#### 2.1.1 Train Test Split

```
[27]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(Poly_x,y,test_size=0.
↪2,random_state=42)
```

### 2.2 4. Modelling and Evaluation

```
[28]: # Modeling with default paramaters
from sklearn.linear_model import LinearRegression
poly_model = LinearRegression()
poly_model.fit(x_train,y_train)

# Prediction
train_predictions = poly_model.predict(x_train)
test_predictions = poly_model.predict(x_test)

# Evaluation
print('Train R2: ',poly_model.score(x_train,y_train))
print('Test R2: ',poly_model.score(x_test,y_test))

from sklearn.model_selection import cross_val_score
print("CV Score: ",cross_val_score(poly_model,x,y,cv=5).mean())
```

Train R2: 0.8369153487751684  
 Test R2: 0.8561470786317726  
 CV Score: 0.7458418619962311

### 2.2.1 Ploy HyperParameter Tuning

```
[29]: train_r2=[]
test_r2=[]
for i in range(1,10):
    from sklearn.preprocessing import PolynomialFeatures
    Poly_converter= PolynomialFeatures(degree=i,include_bias=False)
    Poly_x = pd.DataFrame(Poly_converter.fit_transform(x))

    from sklearn.model_selection import train_test_split
    x_train,x_test,y_train,y_test = train_test_split(Poly_x,y,test_size=0.
↪2,random_state=42)

    from sklearn.linear_model import LinearRegression
    model = LinearRegression()
    model.fit(x_train,y_train)

    test_r2.append(model.score(x_test,y_test))
    train_r2.append(model.score(x_train,y_train))

print(pd.DataFrame([test_r2,train_r2]))
```

	0	1	2	3	4	5	6 \
0	0.768277	0.856147	0.853815	0.847705	0.766139	0.399658	-188.797616
1	0.743605	0.836915	0.843785	0.856941	0.870320	0.881784	0.896520

	7	8
0	-2047.899405	-18944.226103
1	0.904957	0.898268

### 2.2.2 Rebuilding Model with best HPT Parameters

```
[30]: final_Poly_converter= PolynomialFeatures(degree=4,include_bias=False)
Pol_x = final_Poly_converter.fit_transform(x)

x_train,x_test,y_train,y_test = train_test_split(Pol_x,y,test_size=0.
↪2,random_state=42)

final_model = LinearRegression()
final_model.fit(x_train,y_train)

poly_train = final_model.score(x_train,y_train)
poly_test = final_model.score(x_test,y_test)
```

```
poly_cv = cross_val_score(final_model,x,y,cv=5).mean()

print('Train R2: ',poly_train)
print('Test R2: ',poly_test)
print("CV Score: ",poly_cv)
```

Train R2: 0.8569408723742798  
 Test R2: 0.8477046103010456  
 CV Score: 0.7458418619962311

## 3 Lasso Regression

### 3.1 4. Modeling and Evaluation

```
[31]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.
↳2,random_state=42)
```

```
[32]: # Modeling with default paramaters
from sklearn.linear_model import Lasso
lasso_base = Lasso()
lasso_base.fit(x_train,y_train)

# Prediction
train_predictions = lasso_base.predict(x_train)
test_predictions = lasso_base.predict(x_test)

# Evaluation
print('Train R2: ',lasso_base.score(x_train,y_train))
print('Test R2: ',lasso_base.score(x_test,y_test))

from sklearn.model_selection import cross_val_score
print("CVS: ",cross_val_score(lasso_base,x,y,cv=5).mean() )
```

Train R2: 0.7436049167576863  
 Test R2: 0.7682845084591609  
 CVS: 0.7458481242772766

```
[33]: # HyperParameter tuning
from sklearn.model_selection import GridSearchCV

estimator = Lasso()

#param_grid = {"alpha":[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1,10,20,30,50,100]}
param_grid = {"alpha":list(range(1,100))}

model_hp = GridSearchCV(estimator,param_grid,cv =↳
↳5,scoring='neg_mean_squared_error')
```

```
model_hp.fit(x_train,y_train)
model_hp.best_params_
```

[33]: {'alpha': 27}

```
[34]: # Final Modeling with best HyperParameter tuning parameters
from sklearn.linear_model import Lasso
lasso_final = Lasso(alpha=27)
lasso_final.fit(x_train,y_train)

print('Intercepr: ',lasso_final.intercept_)
print('Coefficients: ',lasso_final.coef_)

# Prediction
train_predictions = lasso_final.predict(x_train)
test_predictions = lasso_final.predict(x_test)

# Evaluation

lasso_train = lasso_final.score(x_train,y_train)
lasso_test = lasso_final.score(x_test,y_test)
lasso_cv = cross_val_score(lasso_final,x,y,cv=5).mean()

print('Train R2: ',lasso_train)
print('Test R2: ',lasso_test)
print("CV Score: ",lasso_cv)
```

```
Intercepr:  -11800.259640347304
Coefficients: [ 257.25689231   -0.          313.71245413   496.51744208
 23445.34029312]
Train R2:   0.7435483160182876
Test R2:   0.7684044851038989
CV Score:  0.745964401137262
```

## 4 Ridge Regression

```
[35]: # Modeling with default paramaters
from sklearn.linear_model import Ridge
ridge_base = Ridge()
ridge_base.fit(x_train,y_train)

# Prediction
train_predictions = ridge_base.predict(x_train)
test_predictions = ridge_base.predict(x_test)

# Evaluation
print('Train R2: ',ridge_base.score(x_train,y_train))
```

```
print('Test R2: ',ridge_base.score(x_test,y_test))

from sklearn.model_selection import cross_val_score
print("CV Score: ",cross_val_score(ridge_base,x,y,cv=5).mean())
```

Train R2: 0.7435838281857797  
 Test R2: 0.7682654944819229  
 CV Score: 0.7458250060836376

```
[36]: # HyperParameter tuning
from sklearn.model_selection import GridSearchCV

estimator = Ridge()

param_grid = {"alpha":[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1,10,20,30,50,100]}

model_hp = GridSearchCV(estimator,param_grid,cv = 5,
    ↪scoring='neg_mean_squared_error')
model_hp.fit(x_train,y_train)
model_hp.best_params_
```

[36]: {'alpha': 0.3}

```
[37]: # Rebuilding the final model with best HyperParameter tuning parameters

# Modeling with default parameters
ridge_final = Ridge(alpha=0.3)
ridge_final.fit(x_train,y_train)

# Prediction
train_predictions = ridge_final.predict(x_train)
test_predictions = ridge_final.predict(x_test)

# Evaluation

ridge_train = ridge_final.score(x_train,y_train)
ridge_test = ridge_final.score(x_test,y_test)
ridge_cv = cross_val_score(ridge_final,x,y,cv=5).mean()

print('Train R2: ',ridge_train)
print('Test R2: ',ridge_test)
print("CV Score: ",ridge_cv)
```

Train R2: 0.7436030782325043  
 Test R2: 0.7682781284991226  
 CV Score: 0.7458412533777042



## 5 ElasticNet Regression

```
[38]: # Modeling with default paramaters
from sklearn.linear_model import ElasticNet
ElasticNet_base = ElasticNet()
ElasticNet_base.fit(x_train,y_train)

# Prediction
train_predictions = ElasticNet_base.predict(x_train)
test_predictions = ElasticNet_base.predict(x_test)

# Evaluation
print('Train R2: ',ElasticNet_base.score(x_train,y_train))
print('Test R2: ',ElasticNet_base.score(x_test,y_test))

from sklearn.model_selection import cross_val_score
print("CV Score: ",cross_val_score(ElasticNet_base,x,y,cv=5).mean())
```

```
Train R2:  0.3823176876430222
Test R2:  0.42375809808088793
CV Score:  0.3869970168655764
```

```
[39]: # HyperParameter tuning
from sklearn.model_selection import GridSearchCV

estimator = ElasticNet()

param_grid = {"alpha":list(range(1,100)),
              'l1_ratio':[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1]}

model_hp = GridSearchCV(estimator,param_grid,cv = 5,scoring='neg_mean_squared_error')
model_hp.fit(x_train,y_train)
model_hp.best_params_
```

```
[39]: {'alpha': 27, 'l1_ratio': 1}
```

```
[40]: # Rebuild the final model with best HyperParameter tuning parameters

ElasticNet_best = ElasticNet(alpha=27,l1_ratio=1)
ElasticNet_best.fit(x_train,y_train)

# Prediction
train_predictions = ElasticNet_best.predict(x_train)
test_predictions = ElasticNet_best.predict(x_test)

# Evaluation
```

```

enr_train = ElasticNet_best.score(x_train,y_train)
enr_test = ElasticNet_best.score(x_test,y_test)
enr_cv = cross_val_score(ElasticNet_best,x,y,cv=5).mean()

print('Train R2: ',enr_train)
print('Test R2: ',enr_test)
print("CV Score: ",enr_cv)

```

```

Train R2: 0.7435483160182876
Test R2: 0.7684044851038989
CV Score: 0.745964401137262

```

```

[41]: d1 = {'Algorithms':['Multiple Linear Regression','Polynomial Regression','Lasso_
Regression','Ridge Regression','ElasticNet Regression'],
        'Train R2':[mlr_train,poly_train,lasso_train,ridge_train,enr_train],
        'Test R2':[mlr_test,poly_test,lasso_test,ridge_test,enr_test],
        'CV Scores':[mlr_cv,poly_cv,lasso_cv,ridge_cv,enr_cv],
        'Assumption for SLR and MLR':['No','NaN','NaN','NaN','NaN']}
all_frame = pd.DataFrame(d1)
all_frame

```

```

[41]:
      Algorithms  Train R2  Test R2  CV Scores  \
0  Multiple Linear Regression  0.743605  0.768277  0.745842
1      Polynomial Regression  0.856941  0.847705  0.745842
2          Lasso Regression  0.743548  0.768404  0.745964
3          Ridge Regression  0.743603  0.768278  0.745841
4      ElasticNet Regression  0.743548  0.768404  0.745964

      Assumption for SLR and MLR
0                               No
1                             NaN
2                             NaN
3                             NaN
4                             NaN

```